



BullionEdge Pro

PRE-TRADE ANALYSIS DESK FOR GOLD

Most traders don't lose on analysis. They lose on the **trade they shouldn't have taken.**

BullionEdge Pro is a decision checkpoint that sits between your chart and your broker. It scores a setup against a published rulebook, sizes the position, and — more often than you'd expect — tells you not to take it.

NO TRADE

XAU/USD • M15

Confluence 58/100

The daily is bearish and this is a long, with no shift in structure to justify it.

Reward-to-risk 1.4 — below your minimum of 2.0
CPI lands in 11 minutes

The account is rarely killed by a bad read

It's killed by the four things that happen after it.

Size by feel

The stop is 40 pips one day and 180 the next, but the lot size doesn't move with it. One trade quietly carries five times the risk of another.

Trade the news spike

CPI lands, the spread triples, the first move reverses inside the hour. The setup was fine. The timing was not.

Revenge after a loss

The rules you set calmly on Sunday are the ones you override on Wednesday, at exactly the moment they were written for.

Every one of those is a decision made before the order, not a chart-reading error. That is the gap this product sits in — and the reason it is a checkpoint rather than another indicator.

What it is, and what it deliberately is not

WHAT IT IS

- ◆ A **pre-trade checklist** that scores a setup against rules you can read
- ◆ A **position sizer** that works back from your stop, not your hopes
- ◆ A **risk governor** that shuts the desk when the day's limit is spent
- ◆ A **journal** that turns your own results into feedback


WHAT IT IS NOT

- ◆ Not a **signal service** — it never tells you what will happen
- ◆ Not a **robot** — there is no order path in the product at all
- ◆ Not a **predictor** — the score is a measure of agreement, not probability
- ◆ Not a **black box** — every verdict shows the checks behind it, including the failures

This distinction is the product. A tool that only ever says *yes* is a signal service with extra clicks — and the honest version is worth more to a serious trader than another arrow on a chart.

Seven steps, in the order a trade actually happens

You tick only what you can actually see on your own chart, and the engine weighs it — a step you cannot honestly tick counts against the score.



Nothing here places an order. The output is a decision and a lot size you type into your own terminal.

Buy / Sell candidate

Entry zone, stop, two targets and a lot size — with the reasoning that produced them.

Wait

The idea is sound but the trigger has not printed. Waiting for the close is what stops this becoming a chase.

No trade

Named reasons, not a shrug — and the most common answer on a quiet day.

The timeframes above your entry have to agree

HIGHER TIMEFRAME

The tide — is this with or against the big picture?

DIRECTIONAL

The swing you are actually trading into

EXECUTION

Where the trigger has to print, on a closed candle

Choose any of the seven — D1, H4, H1, M30, M15, M5, M1 — and the two tiers above it move with your choice.

Why that matters

A scalper on M1 and a swing trader on H4 are not doing the same job. The desk reads the tiers **relative to where you actually enter**, so a swing entry is judged against the daily picture — not against the fifteen-minute one.

And when it can't

Entering off the daily, there is nothing above it. The desk **says so** rather than inventing a higher timeframe to agree with — because a confirmation that isn't there is worse than no confirmation.

The feature competitors don't advertise: **it refuses**

NO TRADE

Higher timeframe opposes the entry, with no shift in structure

NO TRADE

Reward-to-risk below your own minimum

STAND DOWN

A high-impact release lands inside fifteen minutes

DESK CLOSED

The daily loss limit is spent — until tomorrow

A tool that can only say yes is a salesman

Every refusal names its reason and shows the checks that produced it — including the ones that passed. The trader can disagree, but they cannot say they weren't told.

Set when calm, enforced when not

The daily loss limit is the feature most worth leaving on. It is chosen on a quiet evening and applied on the afternoon you most want to override it — which is exactly when it is most likely to be right.

The lot size is an output, never an input

$$\begin{array}{|c|} \hline \text{ACCOUNT} \\ \hline \$5,000 \\ \hline \end{array} \times \begin{array}{|c|} \hline \text{RISK / TRADE} \\ \hline 0.5\% \\ \hline \end{array} \div \begin{array}{|c|} \hline \text{STOP DISTANCE} \\ \hline \text{structural} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{POSITION SIZE} \\ \hline 0.12 \text{ lots} \\ \hline \$25 \text{ at risk — the same \$25 on every trade} \\ \hline \end{array}$$

The stop goes where the structure says. The lot size then follows from it — so a wide stop means a smaller position, not a bigger loss.

Broker-aware

Contract size, tick value and lot step come from your broker's spec, so the number is one you can actually type.

Honest about spread

Where the feed publishes candles rather than dealing prices, it says the spread cannot be measured instead of guessing.

One-tap handoff

The finished ticket — pair, direction, entry, stop, targets, lots — copies field by field into your own terminal.

You do not need to be right often. You need **R to mean the same thing every time.**

WIN RATE	AVERAGE WIN	PER TRADE
30%	3.0R	+0.20R
40%	2.0R	+0.20R
50%	1.5R	+0.25R
60%	1.0R	+0.20R
40%	1.0R	-0.20R

1R = the amount risked on the trade. Losses assumed at 1R, which is what a stop that is actually honoured produces. Arithmetic only — not a forecast, and not anybody's results.

The desk does not raise your win rate. Nothing can promise that. What it does is hold the two columns that make the arithmetic work — a constant R, and a reward-to-risk floor you set once and are then held to.

Losing on 6 trades in 10 and still ahead

The 40% row is profitable **only because the average win is 2R**. Drop that to 1R and the identical win rate turns the account over. That single column is what the minimum reward-to-risk setting defends.

And the assumption underneath all of it

Every row assumes **R is constant**. Size by feel — bigger when confident, smaller when shaken — and R stops being a unit. The table stops describing your account, and one oversized loss erases twenty disciplined wins.

Five leaks, and the part of the desk that closes each

THE LEAK	WHAT IT COSTS	WHAT CLOSES IT
Size by feel	R stops being constant, so the expectancy table no longer describes your account	Position sizer — lots derived from the stop, never typed
Trading the news spike	Spread widens, the first move reverses, the stop is hit on noise	Dated gold calendar — stand-down inside 15 minutes
Revenge after a loss	The largest single-day losses in most journals sit here	Daily loss limit — the desk closes until tomorrow
Taking the weak setup	A 40-score entry taken on hope, which the trader already knew was thin	Score written down before entry; skipping costs nothing
No honest record	The same mistake repeats for months because nothing counted it	Journal — outcomes, emotion, and whether you followed the plan

Not one of these is a chart-reading problem — and not one is fixed by a better indicator. They are decisions made in the sixty seconds before the order, which is the only window this product operates in.

A gold calendar, not a generic one

Every date is taken from the agency that publishes it — and every row says **which way a surprise pushes gold**, in both directions.

RELEASE	SOURCE	IMPACT
FOMC rate decision	Federal Reserve	EXTREME
Non-farm payrolls	BLS	EXTREME
CPI — consumer inflation	BLS	EXTREME
PPI, Core PCE, GDP	BLS / BEA	HIGH
Retail sales, ISM PMIs	Census / ISM	HIGH

Anything the app cannot date, it does not date. Michigan sentiment publishes only its next release, so it is listed by time of day and labelled as such rather than guessed at.

Example — CPI

Hot: knocks gold down on the day, because higher-for-longer lifts real yields faster than the hedge bid arrives.

Soft: supports gold — cuts get priced in and the dollar softens.

Why both directions

An impact rating tells a trader a release matters and nothing about what to do when the number lands. Gold competes with real yields and is priced in dollars, so it usually moves **opposite** the headline — which is the part worth stating.

The only part that learns is your own record

- ◆ **Every analysis is logged** — the score, the checks, the plan, and what you actually did
- ◆ **Outcomes close the loop** — win, loss, break-even, in money and in R
- ◆ **Patterns surface from aggregates** — win rate by session, by score band, by emotion
- ◆ **Discipline is measured**, not assumed — how often you overrode your own rules

The coach reads **aggregates only** — never your notes or individual trades. And with a handful of trades it says the numbers are noise rather than dressing them up as an edge.

READ OF YOUR LOG

Across **24** closed trades you are winning **50%** with an expectancy of **0.70R** per trade.

Profit factor 2.45 — your winners are meaningfully outpacing your losers.

Discipline score 55/100. Discipline is the leak here — you are overriding your own rules often enough that the rules are not doing their job.

Illustrative figures, from the app's own wording.

Six stages, each with a number the app checks against your journal

You do not decide which stage you are on. It is measured from what you logged — which is why it is worth something.

1 Record

You cannot improve what you do not measure.

10 trades logged with an outcome

2 Defined risk

Every trade had a stop, and the size respected it.

90% of the last 20 carry a stop and a risk figure

3 Discipline

You followed the plan you wrote, including when it cost you.

80% of the last 20 followed the plan, no revenge trades

4 Selectivity

You stopped taking the setups you already knew were weak.

70% of taken trades scored 55 or better

5 Positive expectancy

Over a real sample, the maths comes out ahead.

30 closed trades and expectancy above zero

6 Consistency

It holds up over time, not just over one good month.

60 closed, positive in 2 of the last 3 blocks of 20, drawdown under 15%

Stage six is passed by **not undoing the previous five**. In the app's own words: most accounts die between stage five and stage six.

Honestly, what the first hundred trades look like

- Week one
You take fewer trades. It feels like a downgrade.
- 10 trades
Stage 1 clears. You have a record, not yet an edge.
- 20 trades
The numbers stop being noise. Patterns by session appear.
- 30 trades
Expectancy starts to mean something.
- 100 trades
The figures settle. Now they describe you.

What it cannot do

It **cannot make a losing strategy profitable**. If the edge is not there, a disciplined process loses money more slowly and more legibly — which is useful, and is not the same thing.

It **cannot predict**. No score here is a probability.

It **cannot stop you overriding it**. It can only make the override a thing you did on purpose, and count it.

The realistic first win

Not a bigger account in week one. It is the week you look at the log and find **no trade you cannot explain** — no oversized loss, no entry taken through a release, no revenge trade after a red morning. That is the state the arithmetic needs before it can work for you.

Built so there is nothing to take

It cannot place a trade

There is no order path in the product. The optional broker link is **read-only** — balance, equity and closed trades — stores no trading password, and never touches your money.

Your record stays local

Trades, notes and settings live in your browser on your device. There is no account database holding your trading history.

Works offline

The full checklist, risk maths and trade plan run with no connection. Aeroplane mode is a fair test of it.

The rulebook is published

Every verdict shows the checks behind it, including the ones that failed. Nothing scores in a black box.

The AI is fenced in

Optional, off by default. Every reply is scanned before you see it: a price the engine did not compute, or any claim to know the future, is flagged above the answer.

Billed by Lemon Squeezy

Merchant of record for payment, invoices and refunds. Card details never reach us.

ACCESS

Try it for a week before anything is charged

\$9.99

WEEK PASS · 7 DAYS

\$31.99

STARTER · 30 DAYS

\$89.99

PRO · 90 DAYS

\$159.99

ELITE · 180 DAYS

7 days free on Starter

Card details taken at signup, **\$0 charged today**, first charge on day 8. Cancel in one click before then and it costs nothing. Your journal and settings stay on your device either way.

Or 10% off any plan

Code **GOLD10** — works on any plan, including after a trial ends. Access is a licence key tied to your device; no account server, nothing to be breached.

Every plan is the whole desk. The tiers are length of access, not features — there is no build with the risk governor removed to sell you a better one. A licence key activates the device you are on; the journal it builds stays yours whether or not you renew.

IN ONE LINE

It is the **pause** between the setup you like and the order you were about to send.

See it live

bullionedgepro.github.io/bullionedge-pro

Runs in the browser. Installs to a phone home screen as an app. Nothing to download.

Honest closing note

This is a decision tool, not financial advice, and it makes no claim about your results. Trading gold is high-risk and you can lose money. What it promises is a process you can inspect — and a rule that stops you on the day you would not have stopped yourself.